

## ▼ Coursework1: Convolutional Neural Networks

### ▼ instructions

Please submit a version of this notebook containing your answers **together with your trained model** on CATe as CW2.zip. Write your answers in the cells below each question.

### ▼ Setting up working environment

For this coursework you will need to train a large network, therefore we recommend you work with Google Colaboratory, which provides free GPU time. You will need a Google account to do so.

Please log in to your account and go to the following page: <https://colab.research.google.com>. Then upload this notebook.

For GPU support, go to "Edit" -> "Notebook Settings", and select "Hardware accelerator" as "GPU".

You will need to install pytorch by running the following cell:

```
1 !pip install torch torchvision

Requirement already satisfied: torch in /usr/local/lib/python3.6/dist-packages (1.7.0+cu101)
Requirement already satisfied: torchvision in /usr/local/lib/python3.6/dist-packages (0.8.1+cu101)
Requirement already satisfied: numpy in /usr/local/lib/python3.6/dist-packages (from torch) (1.19.5)
Requirement already satisfied: typing-extensions in /usr/local/lib/python3.6/dist-packages (from torch) (3.7.4.3)
Requirement already satisfied: future in /usr/local/lib/python3.6/dist-packages (from torch) (0.16.0)
Requirement already satisfied: dataclasses in /usr/local/lib/python3.6/dist-packages (from torch) (0.8)
Requirement already satisfied: pillow>=4.1.1 in /usr/local/lib/python3.6/dist-packages (from torchvision) (7.0.0)
```

### ▼ Introduction

For this coursework you will implement one of the most commonly used model for image recognition tasks, the Residual Network. The architecture is introduced in 2015 by Kaiming He, et al. in the paper ["Deep residual learning for image recognition"](#).

In a residual network, each block contains some convolutional layers, plus "skip" connections, which allow the activations to by pass a layer, and then be summed up with the activations of the skipped layer. The image below illustrates a building block in residual networks.



Depending on the number of building blocks, resnets can have different architectures, for example ResNet-50, ResNet-101 and etc. Here you are required to build ResNet-18 to perform classification on the CIFAR-10 dataset, therefore your network will have the following architecture:



### ▼ Part 1 (40 points)

In this part, you will use basic pytorch operations to define the 2D convolution, max pooling operation, linear layer as well as 2d batch normalization.

### ▼ YOUR TASK

- implement the forward pass for Conv2D, MaxPool2D, Linear and BatchNorm2d
- You are **NOT** allowed to use the torch.nn modules

```
1 import torch
2 import torch.nn as nn
3 import torch.nn.functional as F
4
5
6 torch.manual_seed(0)
7
8 class Conv2d(nn.Module):
9     def __init__(self,
10                 in_channels,
11                 out_channels,
12                 kernel_size,
13                 stride=1,
```

04/02/2021460cw1\_2020.ipynb - Colaboratory

```
14         padding=0,
15         bias=True):
16
17     super(Conv2d, self).__init__()
18     """
19     An implementation of a convolutional layer.
20
21     The input consists of N data points, each with C channels, height H and
22     width W. We convolve each input with F different filters, where each fil
23     spans all C channels and has height HH and width WW.
24
25     Parameters:
26     - w: Filter weights of shape (F, C, HH, WW)
27     - b: Biases, of shape (F,)
28     - kernel_size: Size of the convolving kernel
29     - stride: The number of pixels between adjacent receptive fields in the
30       horizontal and vertical directions.
31     - padding: The number of pixels that will be used to zero-pad the input.
32     """
33
34     #####
35     # TODO: Define the parameters used in the forward pass #
36     #####
37     # *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
38     torch.manual_seed(0)
39     self.C = in_channels
40     self.F = out_channels
41     self.kernel_size = kernel_size
42
43     self.stride = stride
44     self.padding = padding
45
46
47     self.w = nn.Parameter(torch.ones(self.F, self.C, self.kernel_size, self.
48     nn.init.xavier_uniform_(self.w)
49
50
51     if bias:
52         self.bias = nn.Parameter(torch.Tensor(self.F,))
53         nn.init.zeros_(self.bias)
54
55     else:
56         self.bias = None
57
58     # *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
59     #####
60     #                               END OF YOUR CODE #
61     #####
62
63     def forward(self, x):
64         """
65         Input:
66         - x: Input data of shape (N, C, H, W)
67         Output:
68         - out: Output data, of shape (N, F, H', W').
69         """
70
71         #####
72         # TODO: Implement the forward pass #
73         #####
74         # *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
75
76         N = x.shape[0]
77         C = x.shape[1]
78         H = x.shape[2]
79         W = x.shape[3]
80
81         out_H = int((H - self.kernel_size + 2*self.padding) / self.stride) + 1
82
83         out_W = int((W - self.kernel_size + 2*self.padding) / self.stride) + 1
84
85
86
87         # Convolution is equivalent with Unfold + Matrix Multiplication + Fold
88
89
90         unfold = F.unfold(x, self.kernel_size, dilation=1, padding= self.padding
91
92         unfold_trans =  unfold.transpose(1,2)
93
94         unfold_matmul = unfold_trans @ self.w.view(self.F, -1).transpose(0,1)
95
96         output_trans = unfold_matmul.transpose(1,2)
97
98         if self.bias != None:
99             output_trans += self.bias
100
```

https://colab.research.google.com/drive/1WTIjbkwrIhcNPWaLC9V4adQzPrSRvmy9#scrollTo=GTcFtkKH1Jlj&uniqifier=2&printMode=true2/18

```
101     output = output_trans.view(N, self.F, out_H, out_W)
102
103     return output
104
105     # *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
106     #####
107     #                               END OF YOUR CODE                               #
108     #####
109
110     return out
111
112
```

```
1 # TESTS #
2 torch.manual_seed(0)
3 input = torch.randn(1, 1, 5, 5)
4
5 conv_real = nn.Conv2d(1, 1, 3, stride= 1, bias=False)
6 conv_real.weight = torch.nn.Parameter(torch.ones((1, 1, 3, 3)))
7 conv_1 = conv_real(input)
8 print(conv_1)
9
10 torch.manual_seed(0)
11
12 conv = Conv2d(1, 1, 3, stride= 1, bias=False)
13 nn.init.ones_(conv.w)
14 conv_2 = conv.forward(input)
15 print(conv_2)

tensor([[[[-2.3918, -2.6389, -0.8890],
          [ 0.6708,  1.8483,  2.2644],
          [ 1.7742,  0.3874,  0.9557]]]], grad_fn=<ThnnConv2DBackward>)
tensor([[[[-2.3918, -2.6389, -0.8890],
          [ 0.6708,  1.8483,  2.2644],
          [ 1.7742,  0.3874,  0.9557]]]], grad_fn=<ViewBackward>)
```

```
1 class MaxPool2d(nn.Module):
2     def __init__(self, kernel_size):
3         super(MaxPool2d, self).__init__()
4         """
5         An implementation of a max-pooling layer.
6
7         Parameters:
8         - kernel_size: the size of the window to take a max over
9         """
10        #####
11        # TODO: Define the parameters used in the forward pass #
12        #####
13        # *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
14
15        self.kernel_size = kernel_size
16
17        # *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
18        #####
19        #                               END OF YOUR CODE                               #
20        #####
21
22    def forward(self, x):
23        """
24        Input:
25        - x: Input data of shape (N, C, H, W)
26        Output:
27        - out: Output data, of shape (N, F, H', W').
28        """
29        #####
30        # TODO: Implement the forward pass #
31        #####
32        # *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
33
34        # Unfold turns each box of the shape of the filter into a
35        # column of its own which we then transpose
36
37        N = x.shape[0]
38        C = x.shape[1] # C equals output channel size (F)
39        H = x.shape[2]
40        W = x.shape[3]
41
42        H_out = int(H // self.kernel_size)
43        W_out = int(W // self.kernel_size)
44
45
46        unfold_input = F.unfold(x, self.kernel_size, stride=self.kernel_size).tr
47
48        unfold_input = unfold_input.view(N, -1, C, self.kernel_size**2).transpos
49
50        out = torch.max(unfold_input, 3)[0].view(N, C, H_out, W_out)
51
```

```
51
52     return out
53
54     # *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
55     #####
56     #                               END OF YOUR CODE                               #
57     #####
58
59     return out


1 # TESTS #
2 torch.manual_seed(10)
3 max_me = MaxPool2d(2)
4 max_real = nn.MaxPool2d(2)
5
6 # Tried with three different inputs to make sure they work
7 #input1 = torch.Tensor([[[[6,1,2,4,7,8],[1,6,7,8,9,3], [3,5,2,0,2,5], [1,2,3,4,3
8 #input2 = torch.Tensor([[[[6,1,2,4,7],[1,6,7,8,9], [3,5,2,0,2], [1,2,3,4,3],[6,1
9 input3 = torch.randn(1, 2, 6, 6)
10
11 ans_me = max_me.forward(input3)
12 print(ans_me)
13 ans_real = max_real.forward(input3)
14 print(ans_real)


        tensor([[[[ 2.6278,  1.0051,  0.4765],
                    [ 1.1167,  1.0763,  1.1209],
                    [ 0.2456,  1.3290,  0.2121]],
                  [[ 0.9741,  0.1490, -0.8832],
                    [-0.4811,  2.7470,  1.3373],
                    [ 0.0948,  1.2828,  1.3399]]]])
        tensor([[[[ 2.6278,  1.0051,  0.4765],
                    [ 1.1167,  1.0763,  1.1209],
                    [ 0.2456,  1.3290,  0.2121]],
                  [[ 0.9741,  0.1490, -0.8832],
                    [-0.4811,  2.7470,  1.3373],
                    [ 0.0948,  1.2828,  1.3399]]]])


1 class Linear(nn.Module):
2     def __init__(self, in_channels, out_channels, bias=True):
3         super(Linear, self).__init__()
4         """
5         An implementation of a Linear layer.
6
7         Parameters:
8         - weight: the learnable weights of the module of shape (in_channels, out
9         - bias: the learnable bias of the module of shape (out_channels).
10        """
11        #####
12        # TODO: Define the parameters used in the forward pass #
13        #####
14        # *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
15
16        self.in_channels = in_channels
17        self.out_channels = out_channels
18
19
20        self.w = nn.Parameter(torch.ones(self.in_channels, self.out_channels))
21        torch.nn.init.xavier_uniform_(self.w)
22
23        if bias:
24            self.bias = nn.Parameter(torch.zeros(self.out_channels,))
25            torch.nn.init.zeros_(self.bias)
26
27        else:
28            self.bias = None
29
30
31
32        # *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
33        #####
34        #                               END OF YOUR CODE                               #
35        #####
36
37    def forward(self, x):
38        """
39        Input:
40        - x: Input data of shape (N, *, H) where * means any number of additiona
41        dimensions and H = in_channels
42        Output:
43        - out: Output data of shape (N, *, H') where * means any number of addit
44        dimensions and H' = out_channels
45        """
46        #####
47        # TODO: Implement the forward pass #
48        #####
```

04/02/2021460cw1\_2020.ipynb - Colaboratory

```
48 #####
49 # *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
50
51 N = x.shape[0]
52 H = x.shape[-1]
53
54 output = x @ self.w
55
56 if self.bias != None:
57     output += self.bias
58
59 return output
60
61 # *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
62 #####
63 #                                     END OF YOUR CODE                                     #
64 #####
65
66 return out


1 # TESTS #
2 torch.manual_seed(10)
3 input = torch.Tensor([3, 5, 4])
4
5 linear = Linear(3, 4, bias=False)
6 nn.init.ones_(linear.w)
7 linear_me = linear.forward(input)
8 print(linear_me)
9
10 linear_torch = nn.Linear(3, 4, bias=False)
11 # the 4 and the 3 are the other way around here as
12 # pytorch module defines weight as
13 # self.weight = Parameter(torch.Tensor(out_features, in_features))
14 linear_torch.weight = nn.Parameter(torch.ones(4, 3))
15 linear_t = linear_torch.forward(input)
16 print(linear_t)


    tensor([12., 12., 12., 12.], grad_fn=<SqueezeBackward3>)
    tensor([12., 12., 12., 12.], grad_fn=<SqueezeBackward3>)


1 class BatchNorm2d(nn.Module):
2     def __init__(self, num_features, eps=1e-05, momentum=0.1, training= True):
3         super(BatchNorm2d, self).__init__()
4         """
5         An implementation of a Batch Normalization over a mini-batch of 2D input
6
7         The mean and standard-deviation are calculated per-dimension over the
8         mini-batches and gamma and beta are learnable parameter vectors of
9         size num_features.
10
11         Parameters:
12         - num_features: C from an expected input of size (N, C, H, W).
13         - eps: a value added to the denominator for numerical stability. Default
14         - momentum: momentum – the value used for the running_mean and running_v
15         computation. Default: 0.1
16         - gamma: the learnable weights of shape (num_features).
17         - beta: the learnable bias of the module of shape (num_features).
18         """
19         #####
20         # TODO: Define the parameters used in the forward pass                                     #
21         #####
22         # *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
23         torch.manual_seed(10)
24         self.num_features = num_features
25         self.eps = eps
26         self.momentum = momentum
27
28         self.gamma = nn.Parameter(torch.ones(1, self.num_features, 1, 1))
29
30         self.beta = nn.Parameter(torch.zeros(1, self.num_features, 1, 1))
31
32         self.moving_mean = torch.zeros(1, self.num_features, 1, 1)
33         self.moving_std = torch.ones(1, self.num_features, 1, 1)
34
35         self.training = training
36
37
38         # *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
39         #####
40         #                                     END OF YOUR CODE                                     #
41         #####
42
43     def forward(self, x):
44         """
45         During training this layer keeps running estimates of its computed mean
46         variance, which are then used for normalization during evaluation.
47         """
```

https://colab.research.google.com/drive/1WTljbkwrhlcNPWaLC9V4adQzPrSRvmy9#scrollTo=GTcFtkKH1Jlj&uniqifier=2&printMode=true5/18



```

46     self.moving_mu = torch.zeros(1, shape[1],
47     self.moving_sigma = torch.ones(1, shape[1])
48
49     if len(shape) == 4:
50         self.moving_mu = torch.zeros(1, shape[1], 1, 1)
51         self.moving_sigma = torch.ones(1, shape[1], 1, 1)
52
53
54     def forward(self, x):
55         # Test
56         if not torch.is_grad_enabled():
57             # PART 3: Test time normalization operation; use self.eps as epsilon
58             # UPDATE:
59             x_hat = (x - self.moving_mu) / torch.sqrt(self.moving_sigma + self.e
60
61             # Logging code for tests; ignore:
62             self._tmp_x_hat_test = x_hat
63
64         # Training
65         else:
66             if len(x.shape) == 2:
67                 # PART 4: Compute mean and var for FC input (retaining feature d
68                 # UPDATE:
69                 mean = x.mean(dim=(0), keepdim=True)
70                 std = x.std(dim=(0), keepdim=True)
71                 var = std**2
72
73             elif len(x.shape) == 4:
74                 # PART 5: Compute mean and var for Conv input (retaining channel
75                 # UPDATE (hint: use `keepdim` flag to use broadcasting later):
76                 mean = x.mean(dim=(0,2,3), keepdim=True)
77                 std = x.std(dim=(0,2,3), keepdim=True)
78                 var = std**2
79             else:
80                 raise ValueError("Incorrect input shape!")
81
82             # Logging code for tests; ignore:
83             self._tmp_mean = mean
84             self._tmp_var = var
85
86             # PART 6: Training time normalization operation; use self.eps as eps
87             # UPDATE:
88             x_hat = (x - mean) / torch.sqrt(var + self.eps)
89
90             # Logging code for tests; ignore:
91             self._tmp_x_hat_train = x_hat
92
93             # PART 7: Updating moving averages; use self.momentum to calculate
94             # contribution to update (hint: be careful about unnecessary
95             # autograd computation tracking)
96             # UPDATE:
97             self.moving_mu = self.momentum * self.moving_mu + (1 - self.momentum
98             self.moving_sigma = self.momentum * self.moving_sigma + (1 - self.mc
99
100             # Logging code for tests; ignore:
101             self._tmp_moving_mu = self.moving_mu
102             self._tmp_moving_sigma = self.moving_sigma
103
104             # PART 8: Scale and shift x_hat using learnable parameters to compute ou
105             # UPDATE:
106             z = x_hat * self.gamma + self.beta
107
108             return z

```

```

1 # TESTS #
2 torch.manual_seed(10)
3
4 input = torch.randn(1, 2, 4, 4)
5 bn = BatchNorm2d(2, training=True)
6 print(bn.forward(input))
7
8 b = BatchNorm((1, 2, 4, 4))
9 print(b.forward(input))
10

```

```

tensor([[[[[-0.5240, -0.2591, -0.5335, -1.0094],
          [ 0.1034, -0.6728, -0.2163,  2.9627],
          [-0.4523,  1.3204,  0.0433,  0.7854],
          [-0.3700, -0.0638, -1.1647,  0.0506]],
          [[ 0.8100,  1.3459,  1.1231,  1.3381],
          [ 1.2637, -0.8392, -0.4752, -0.4079],
          [-0.6139, -0.2653,  1.0377, -0.0633],
          [-0.3269, -1.3390, -1.3681, -1.2196]]]], grad_fn=<AddBackward0>)]
tensor([[[[[-0.5240, -0.2591, -0.5335, -1.0094],
          [ 0.1034, -0.6728, -0.2163,  2.9627],
          [-0.4523,  1.3204,  0.0433,  0.7854],
          [-0.3700, -0.0638, -1.1647,  0.0506]],

```

```
[[ 0.8100,  1.3459,  1.1231,  1.3381],
 [ 1.2637, -0.8392, -0.4752, -0.4079],
 [-0.6139, -0.2653,  1.0377, -0.0633],
 [-0.3269, -1.3390, -1.3681, -1.2196]]]], grad_fn=<AddBackward0>)
```

▼ Part 2

In this part, you will train a ResNet-18 defined on the CIFAR-10 dataset. Code for training and evaluation are provided.

▼ Your Task

- 1. Train your network to achieve the best possible test set accuracy after a maximum of 10 epochs of training.
- 2. You can use techniques such as optimal hyper-parameter searching, data pre-processing
- 3. If necessary, you can also use another optimizer
- 4. **Answer the following question:** Given such a network with a large number of trainable parameters, and a training set of a large number of data, what do you think is the best strategy for hyperparameter searching?

**Q: Given such a network with a large number of trainable parameters, and a training set of a large number of data, what do you think is the best strategy for hyperparameter searching?**

Deep learning models have countless trainable parameters and finding the optimal configuration is not a trivial challenge. There are several potential strategies, Trial and Error, Grid Search and Random Search. Trial and Error is fairly self explanatory, but not particularly effective with such a large number of parameters and a large dataset. With Grid search we just try every possible configuration, but this is a very naive approach and is very not efficient. Random search on the other hand is more effective than Grid search when dealing with a vast number of hyperparameters and a large dataset, and it generally gives better overall results after a lower number of iterations.

The best strategy for hyperparameter searching overall is Bayesian Optimization. This search strategy builds a probability (surrogate) model of the objective function and uses it to select hyperparameter to evaluate in the true objective function. Essentially it tries to predict the metrics we care about from the hyperparameters configuration. As we move through the iterations the probability model becomes more and more confident about which guess leads to improvements. We use the Gaussian process as the probability model that will learn the mapping from hyperparameters configuration to the metric of interest.

```
1 import torch
2 from torch.nn import Conv2d, MaxPool2d
3 import torch.nn as nn
4 import torch.nn.functional as F
5
6 from itertools import product
7
```

Next, we define ResNet-18:

```
1 # define resnet building blocks
2
3 class ResidualBlock(nn.Module):
4     def __init__(self, inchannel, outchannel, stride=1):
5
6         super(ResidualBlock, self).__init__()
7
8         self.left = nn.Sequential(Conv2d(inchannel, outchannel, kernel_size=3,
9                                           stride=stride, padding=1, bias=False),
10                                   nn.BatchNorm2d(outchannel),
11                                   nn.ReLU(inplace=True),
12                                   Conv2d(outchannel, outchannel, kernel_size=3,
13                                         stride=1, padding=1, bias=False),
14                                   nn.BatchNorm2d(outchannel))
15
16         self.shortcut = nn.Sequential()
17
18         if stride != 1 or inchannel != outchannel:
19
20             self.shortcut = nn.Sequential(Conv2d(inchannel, outchannel,
21                                                   kernel_size=1, stride=stride,
22                                                   padding = 0, bias=False),
23                                           nn.BatchNorm2d(outchannel) )
24
25     def forward(self, x):
26
27         out = self.left(x)
28
29         out += self.shortcut(x)
30
31         out = F.relu(out)
```



```
31         out = torch.nn.ReLU(out),
32
33         return out
34
35
36
37     # define resnet
38
39 class ResNet(nn.Module):
40
41     def __init__(self, ResidualBlock, num_classes = 10):
42
43         super(ResNet, self).__init__()
44
45         self.inchannel = 64
46         self.conv1 = nn.Sequential(Conv2d(3, 64, kernel_size = 3, stride = 1,
47                                         padding = 1, bias = False),
48                                     nn.BatchNorm2d(64),
49                                     nn.ReLU())
50
51         self.layer1 = self.make_layer(ResidualBlock, 64, 2, stride = 1)
52         self.layer2 = self.make_layer(ResidualBlock, 128, 2, stride = 2)
53         self.layer3 = self.make_layer(ResidualBlock, 256, 2, stride = 2)
54         self.layer4 = self.make_layer(ResidualBlock, 512, 2, stride = 2)
55         self.maxpool = MaxPool2d(4)
56         self.fc = nn.Linear(512, num_classes)
57
58
59     def make_layer(self, block, channels, num_blocks, stride):
60
61         strides = [stride] + [1] * (num_blocks - 1)
62
63         layers = []
64
65         for stride in strides:
66
67             layers.append(block(self.inchannel, channels, stride))
68
69             self.inchannel = channels
70
71         return nn.Sequential(*layers)
72
73
74     def forward(self, x):
75
76         x = self.conv1(x)
77
78         x = self.layer1(x)
79         x = self.layer2(x)
80         x = self.layer3(x)
81         x = self.layer4(x)
82
83         x = self.maxpool(x)
84
85         x = x.view(x.size(0), -1)
86
87         x = self.fc(x)
88
89         return x
90
91
92 def ResNet18():
93     return ResNet(ResidualBlock)
```

▼ Loading dataset

We will import images from the [torchvision.datasets](#) library

First, we need to define the alterations (transforms) we want to perform to our images - given that transformations are applied when importing the data.

Define the following transforms using the torchvision.datasets library -- you can read the transforms documentation [here](#):

- 1. Convert images to tensor
- 2. Normalize mean and std of images with values:mean=[0.4914, 0.4822, 0.4465], std=[0.2023, 0.1994, 0.2010]

```
1 import torch.optim as optim
2 from torch.utils.data import DataLoader
3 from torch.utils.data import sampler
4
5
6 import torchvision.datasets as dset
7
8 import numpy as np
9
10 import torchvision.transforms as T
```

04/02/2021460cw1\_2020.ipynb - Colaboratory

```
11
12 import torchvision
13 import matplotlib.pyplot as plt
14
15
16
17 torch.set_printoptions(linewidth=120)
18 torch.set_grad_enabled(True)
19 from torch.utils.tensorboard import SummaryWriter
20
21 #####
22 #                               YOUR CODE HERE                               #
23 #####
24
25 transform_normalise = T.Compose([T.ToTensor(),
26                                  T.Normalize((0.4914, 0.4822, 0.4465), (0.2023, 0
27                                  T.RandomHorizontalFlip(),
28                                  T.RandomResizedCrop(32)])
29
30                                  #T.ColorJitter(saturation=0.01, hue=0.01)
31
32 """ColorJitter, RandomHorizontal Flip and RandomResizedCrop were added
33 to increase the robustenss of the training through data augmentation. Tried all
34 and found that the ColorJitter poorly affected our accuracy after 10
35 iterations and therefore left this one out in the final implementation """
36
37
38 #####
39 #                               END OF YOUR CODE                               #
40 #####
41
42
43
44
45
46
47
48
49
50
51
52
53
54
55
56
57
58
59
60
61
62
63
64
65
66
67
68
69
70
71
72
73
74
75
76
77
78
79
80
81
82
83
84
85
86
87
88
89
90
91
92
93
94
95
96
97
98
99
100
101
102
103
104
105
106
107
108
109
110
111
112
113
114
115
116
117
118
119
120
121
122
123
124
125
126
127
128
129
130
131
132
133
134
135
136
137
138
139
140
141
142
143
144
145
146
147
148
149
150
151
152
153
154
155
156
157
158
159
160
161
162
163
164
165
166
167
168
169
170
171
172
173
174
175
176
177
178
179
180
181
182
183
184
185
186
187
188
189
190
191
192
193
194
195
196
197
198
199
200
201
202
203
204
205
206
207
208
209
210
211
212
213
214
215
216
217
218
219
220
221
222
223
224
225
226
227
228
229
230
231
232
233
234
235
236
237
238
239
240
241
242
243
244
245
246
247
248
249
250
251
252
253
254
255
256
257
258
259
260
261
262
263
264
265
266
267
268
269
270
271
272
273
274
275
276
277
278
279
280
281
282
283
284
285
286
287
288
289
290
291
292
293
294
295
296
297
298
299
300
301
302
303
304
305
306
307
308
309
310
311
312
313
314
315
316
317
318
319
320
321
322
323
324
325
326
327
328
329
330
331
332
333
334
335
336
337
338
339
340
341
342
343
344
345
346
347
348
349
350
351
352
353
354
355
356
357
358
359
360
361
362
363
364
365
366
367
368
369
370
371
372
373
374
375
376
377
378
379
380
381
382
383
384
385
386
387
388
389
390
391
392
393
394
395
396
397
398
399
400
401
402
403
404
405
406
407
408
409
410
411
412
413
414
415
416
417
418
419
420
421
422
423
424
425
426
427
428
429
430
431
432
433
434
435
436
437
438
439
440
441
442
443
444
445
446
447
448
449
450
451
452
453
454
455
456
457
458
459
460
461
462
463
464
465
466
467
468
469
470
471
472
473
474
475
476
477
478
479
480
481
482
483
484
485
486
487
488
489
490
491
492
493
494
495
496
497
498
499
500
501
502
503
504
505
506
507
508
509
510
511
512
513
514
515
516
517
518
519
520
521
522
523
524
525
526
527
528
529
530
531
532
533
534
535
536
537
538
539
540
541
542
543
544
545
546
547
548
549
550
551
552
553
554
555
556
557
558
559
560
561
562
563
564
565
566
567
568
569
570
571
572
573
574
575
576
577
578
579
580
581
582
583
584
585
586
587
588
589
590
591
592
593
594
595
596
597
598
599
600
601
602
603
604
605
606
607
608
609
610
611
612
613
614
615
616
617
618
619
620
621
622
623
624
625
626
627
628
629
630
631
632
633
634
635
636
637
638
639
640
641
642
643
644
645
646
647
648
649
650
651
652
653
654
655
656
657
658
659
660
661
662
663
664
665
666
667
668
669
670
671
672
673
674
675
676
677
678
679
680
681
682
683
684
685
686
687
688
689
690
691
692
693
694
695
696
697
698
699
700
701
702
703
704
705
706
707
708
709
710
711
712
713
714
715
716
717
718
719
720
721
722
723
724
725
726
727
728
729
730
731
732
733
734
735
736
737
738
739
740
741
742
743
744
745
746
747
748
749
750
751
752
753
754
755
756
757
758
759
760
761
762
763
764
765
766
767
768
769
770
771
772
773
774
775
776
777
778
779
780
781
782
783
784
785
786
787
788
789
790
791
792
793
794
795
796
797
798
799
800
801
802
803
804
805
806
807
808
809
810
811
812
813
814
815
816
817
818
819
820
821
822
823
824
825
826
827
828
829
830
831
832
833
834
835
836
837
838
839
840
841
842
843
844
845
846
847
848
849
850
851
852
853
854
855
856
857
858
859
860
861
862
863
864
865
866
867
868
869
870
871
872
873
874
875
876
877
878
879
880
881
882
883
884
885
886
887
888
889
890
891
892
893
894
895
896
897
898
899
900
901
902
903
904
905
906
907
908
909
910
911
912
913
914
915
916
917
918
919
920
921
922
923
924
925
926
927
928
929
930
931
932
933
934
935
936
937
938
939
940
941
942
943
944
945
946
947
948
949
950
951
952
953
954
955
956
957
958
959
960
961
962
963
964
965
966
967
968
969
970
971
972
973
974
975
976
977
978
979
980
981
982
983
984
985
986
987
988
989
990
991
992
993
994
995
996
997
998
999
1000
```

'added ColorJitter, RandomHorizontal Flip and RandomResizedCrop to increase the robustenss\nof the training. Tried all three and found that the ColorJitter poorly affected our \naccuracy after 10 iterations and therefore left this one out in the final implementation '

Now load the dataset using the transform you defined above, with batch\_size = 64

You can check the documentation [here](#). Then create data loaders (using DataLoader from torch.utils.data) for the training and test set

```
1
2 #####
3 #                               YOUR CODE HERE                               #
4 #####
5
6 val_percent = 0.1
7
8 batch_size = 64
9
10 data_dir = './data'
11
12 trainval_data = dset.CIFAR10(
13     root = data_dir,
14     train = True,
15     download= True,
16     transform = transform_normalise
17 )
18
19 test_data = dset.CIFAR10(
20     root = data_dir,
21     train = False,
22     download= True,
23     transform = transform_normalise
24 )
25
26
27
28 trainval_length = len(trainval_data)
29
30 val_split_num = int(np.floor(val_percent * trainval_length))
31
32 print(val_split_num)
33
34
35 loader_val = DataLoader(trainval_data, batch_size= batch_size,
36                          sampler= sampler.SubsetRandomSampler(range(val_split_num
37
38 loader_train = DataLoader(trainval_data, batch_size= batch_size,
39                           sampler= sampler.SubsetRandomSampler(range(val_split_n
40
41 print(len(loader_train))
42
43 loader_test = DataLoader(test_data)
44
45
46
47
48
49
50
51
52
53
54
55
56
57
58
59
60
61
62
63
64
65
66
67
68
69
70
71
72
73
74
75
76
77
78
79
80
81
82
83
84
85
86
87
88
89
90
91
92
93
94
95
96
97
98
99
100
101
102
103
104
105
106
107
108
109
110
111
112
113
114
115
116
117
118
119
120
121
122
123
124
125
126
127
128
129
130
131
132
133
134
135
136
137
138
139
140
141
142
143
144
145
146
147
148
149
150
151
152
153
154
155
156
157
158
159
160
161
162
163
164
165
166
167
168
169
170
171
172
173
174
175
176
177
178
179
180
181
182
183
184
185
186
187
188
189
190
191
192
193
194
195
196
197
198
199
200
201
202
203
204
205
206
207
208
209
210
211
212
213
214
215
216
217
218
219
220
221
222
223
224
225
226
227
228
229
230
231
232
233
234
235
236
237
238
239
240
241
242
243
244
245
246
247
248
249
250
251
252
253
254
255
256
257
258
259
260
261
262
263
264
265
266
267
268
269
270
271
272
273
274
275
276
277
278
279
280
281
282
283
284
285
286
287
288
289
290
291
292
293
294
295
296
297
298
299
300
301
302
303
304
305
306
307
308
309
310
311
312
313
314
315
316
317
318
319
320
321
322
323
324
325
326
327
328
329
330
331
332
333
334
335
336
337
338
339
340
341
342
343
344
345
346
347
348
349
350
351
352
353
354
355
356
357
358
359
360
361
362
363
364
365
366
367
368
369
370
371
372
373
374
375
376
377
378
379
380
381
382
383
384
385
386
387
388
389
390
391
392
393
394
395
396
397
398
399
400
401
402
403
404
405
406
407
408
409
410
411
412
413
414
415
416
417
418
419
420
421
422
423
424
425
426
427
428
429
430
431
432
433
434
435
436
437
438
439
440
441
442
443
444
445
446
447
448
449
450
451
452
453
454
455
456
457
458
459
460
461
462
463
464
465
466
467
468
469
470
471
472
473
474
475
476
477
478
479
480
481
482
483
484
485
486
487
488
489
490
491
492
493
494
495
496
497
498
499
500
501
502
503
504
505
506
507
508
509
510
511
512
513
514
515
516
517
518
519
520
521
522
523
524
525
526
527
528
529
530
531
532
533
534
535
536
537
538
539
540
541
542
543
544
545
546
547
548
549
550
551
552
553
554
555
556
557
558
559
560
561
562
563
564
565
566
567
568
569
570
571
572
573
574
575
576
577
578
579
580
581
582
583
584
585
586
587
588
589
590
591
592
593
594
595
596
597
598
599
600
601
602
603
604
605
606
607
608
609
610
611
612
613
614
615
616
617
618
619
620
621
622
623
624
625
626
627
628
629
630
631
632
633
634
635
636
637
638
639
640
641
642
643
644
645
646
647
648
649
650
651
652
653
654
655
656
657
658
659
660
661
662
663
664
665
666
667
668
669
670
671
672
673
674
675
676
677
678
679
680
681
682
683
684
685
686
687
688
689
690
691
692
693
694
695
696
697
698
699
700
701
702
703
704
705
706
707
708
709
710
711
712
713
714
715
716
717
718
719
720
721
722
723
724
725
726
727
728
729
730
731
732
733
734
735
736
737
738
739
740
741
742
743
744
745
746
747
748
749
750
751
752
753
754
755
756
757
758
759
760
761
762
763
764
765
766
767
768
769
770
771
772
773
774
775
776
777
778
779
780
781
782
783
784
785
786
787
788
789
790
791
792
793
794
795
796
797
798
799
800
801
802
803
804
805
806
807
808
809
810
811
812
813
814
815
816
817
818
819
820
821
822
823
824
825
826
827
828
829
830
831
832
833
834
835
836
837
838
839
840
841
842
843
844
845
846
847
848
849
850
851
852
853
854
855
856
857
858
859
860
861
862
863
864
865
866
867
868
869
870
871
872
873
874
875
876
877
878
879
880
881
882
883
884
885
886
887
888
889
890
891
892
893
894
895
896
897
898
899
900
901
902
903
904
905
906
907
908
909
910
911
912
913
914
915
916
917
918
919
920
921
922
923
924
925
926
927
928
929
930
931
932
933
934
935
936
937
938
939
940
941
942
943
944
945
946
947
948
949
950
951
952
953
954
955
956
957
958
959
960
961
962
963
964
965
966
967
968
969
970
971
972
973
974
975
976
977
978
979
980
981
982
983
984
985
986
987
988
989
990
991
992
993
994
995
996
997
998
999
1000
```

https://colab.research.google.com/drive/1WTIjbkwrIhcNPWaLC9V4adQzPrSRvmy9#scrollTo=GTcFtkKH1Jlj&uniqifier=2&printMode=true

```
45
46
47 batch = next(iter(loader_train))
48 len(batch)
49 images, labels = batch
50
51 print(images.shape)
52
53
54 grid = torchvision.utils.make_grid(images)
55 plt.figure(figsize=(15,15))
56 plt.imshow(np.transpose(grid, (1,2,0)))
57
58 print('labels', labels)
59
60 #####
61 #                               END OF YOUR CODE                               #
62 #####
63
64
```

Files already downloaded and verified  
Files already downloaded and verified  
Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).  
5000  
704  
torch.Size([64, 3, 32, 32])  
labels tensor([4, 0, 1, 8, 8, 5, 4, 4, 0, 7, 9, 6, 2, 6, 9, 0, 3, 2, 1, 8, 0, 7, 1, 3, 9, 5, 6, 6, 1, 4, 1, 1, 6, 4, 5, 3, 6, 5, 5, 5, 6, 8, 5, 5, 6, 7, 9, 0, 2, 2, 7, 4, 3, 2, 1, 8, 5, 4, 7, 6, 9, 9])



```
1 USE_GPU = True
2 dtype = torch.float32
3
4 if USE_GPU and torch.cuda.is_available():
5     device = torch.device('cuda')
6 else:
7     device = torch.device('cpu')
8
9
10
11 print_every = 100
12 def check_accuracy(loader, model):
13     # function for test accuracy on validation and test set
14
15     if loader.dataset.train:
16         print('Checking accuracy on validation set')
17     else:
18         print('Checking accuracy on test set')
19     num_correct = 0
```

```

20     num_samples = 0
21     model.eval() # set model to evaluation mode
22     with torch.no_grad():
23         for x, y in loader:
24             x = x.to(device=device, dtype=dtype) # move to device
25             y = y.to(device=device, dtype=torch.long)
26             scores = model(x)
27             _, preds = scores.max(1)
28             num_correct += (preds == y).sum()
29             num_samples += preds.size(0)
30     acc = float(num_correct) / num_samples
31     print('Got %d / %d correct (%.2f)' % (num_correct, num_samples, 100 * ac
32
33     return acc
34
35
36 def train_part(model, optimizer, epochs=1):
37     """
38     Train a model on CIFAR-10 using the PyTorch Module API.
39
40     Inputs:
41     - model: A PyTorch Module giving the model to train.
42     - optimizer: An Optimizer object we will use to train the model
43     - epochs: (Optional) A Python integer giving the number of epochs to train f
44
45     Returns: Nothing, but prints model accuracies during training.
46     """
47     model = model.to(device=device) # move the model parameters to CPU/GPU
48     for e in range(epochs):
49         print(len(loader_train))
50         for t, (x, y) in enumerate(loader_train):
51             model.train() # put model to training mode
52             x = x.to(device=device, dtype=dtype) # move to device, e.g. GPU
53             y = y.to(device=device, dtype=torch.long)
54
55             scores = model(x)
56             loss = F.cross_entropy(scores, y)
57
58             # Zero out all of the gradients for the variables which the optimize
59             # will update.
60             optimizer.zero_grad()
61
62             loss.backward()
63
64             # Update the parameters of the model using the gradients
65             optimizer.step()
66
67             if t % print_every == 0:
68                 print('Epoch: %d, Iteration %d, loss = %.4f' % (e, t, loss.item(
69                     acc = check_accuracy(loader_val, model)
70                     writer.add_scalar('accuracy', acc, e)
71
72                 # added this to add variables to tensorboard to see the impact c
73                 # changing parameters
74                 for name, weight in model.named_parameters():
75                     writer.add_histogram(name, weight, e)
76                     writer.add_histogram(f'{name}.grad', weight.grad, e)
77                 #writer.add_histogram('conv1.weight', model.conv1[0].weight, e)
78                 print()

```

```

1 # code for optimising your network performance
2
3 #####
4 #                               YOUR CODE HERE                               #
5 #####
6 """
7 Learning rate is an optimization parameter that tends to minimise
8 the loss function thereby reduces the model's error. The adam optimizer
9 was intially analysed for a few different learning rates of 0.01, 0.005,
10 0.001 and 0.0005 to get some ideas of good values for the learning
11 rate parameter. Following these results it was found that learning rates
12 around 0.001 were good and therefore a larger experiment was run with
13 many more values focussed around this area to pin point the best
14 learning rate.
15 (graphs are plotted below in Tensorboard for the accuracies of different
16 learning rates)
17
18 """
19 torch.manual_seed(10)
20
21 #parameters = dict(
22 #    lr = [0.005, 0.002, 0.0015, 0.001, 0.0005, 0.0001, 0.00005, 0.00001, 0.0000
23
24 #param_values = [v for v in parameters.values()]
25
26 #for lr in product(*param_values):
27 #    comment = f'lr={lr}'

```

```
27 # comment - 1 11-111}
28
29 # writer is used to send data to tensorboard
30 writer = SummaryWriter()
31
32 #####
33 #                               END OF YOUR CODE                               #
34 #####
35 torch.manual_seed(4)
36
37 # define and train the network
38 model = ResNet18()
39 optimizer = optim.Adam(model.parameters(), lr= 0.001)
40
41 train_part(model, optimizer, epochs = 10)
42
43 # report test set accuracy
44
45 check_accuracy(loader_test, model)
46
47 writer.close()
48 # save the model
49 torch.save(model.state_dict(), 'model2.pt')
```

Epoch: 8, Iteration 200, loss = 0.1278  
Checking accuracy on validation set  
Got 4327 / 5000 correct (86.54)

Epoch: 8, Iteration 300, loss = 0.1644  
Checking accuracy on validation set  
Got 4295 / 5000 correct (85.90)

Epoch: 8, Iteration 400, loss = 0.3116  
Checking accuracy on validation set  
Got 4164 / 5000 correct (83.28)

Epoch: 8, Iteration 500, loss = 0.3613  
Checking accuracy on validation set  
Got 4270 / 5000 correct (85.40)

Epoch: 8, Iteration 600, loss = 0.1884  
Checking accuracy on validation set  
Got 4287 / 5000 correct (85.74)

Epoch: 8, Iteration 700, loss = 0.2238  
Checking accuracy on validation set  
Got 4276 / 5000 correct (85.52)

704  
Epoch: 9, Iteration 0, loss = 0.0855  
Checking accuracy on validation set  
Got 4284 / 5000 correct (85.68)

Epoch: 9, Iteration 100, loss = 0.1088  
Checking accuracy on validation set  
Got 4291 / 5000 correct (85.82)

Epoch: 9, Iteration 200, loss = 0.2386  
Checking accuracy on validation set  
Got 4327 / 5000 correct (86.54)

Epoch: 9, Iteration 300, loss = 0.1937  
Checking accuracy on validation set  
Got 4273 / 5000 correct (85.46)

Epoch: 9, Iteration 400, loss = 0.4077  
Checking accuracy on validation set  
Got 4285 / 5000 correct (85.70)

Epoch: 9, Iteration 500, loss = 0.1939  
Checking accuracy on validation set  
Got 4303 / 5000 correct (86.06)

Epoch: 9, Iteration 600, loss = 0.1228  
Checking accuracy on validation set  
Got 4336 / 5000 correct (86.72)

Epoch: 9, Iteration 700, loss = 0.2030  
Checking accuracy on validation set  
Got 4292 / 5000 correct (85.84)

Checking accuracy on test set  
Got 8339 / 10000 correct (83.39)

1  
2 grid = torchvision.utils.make\_grid(images)  
3  
4 images, labels = images.cuda(), labels.cuda()  
5  
6 #writer.add\_image(tag= 'images', img\_tensor= grid)  
7 #writer.add\_graph(model, images)  
8 #writer.close()



```
1 # used tensorboard to visualise the results
2
3 %load_ext tensorboard
4 %reload_ext tensorboard
5 %tensorboard --logdir=runs
6
```

Reusing TensorBoard on port 6006 (pid 389), started 2:48:34 ago. (Use '!kill 389' to kill it.)

TensorBoard

SCALARS

DISTRIBUTIONS

HISTOGRAMS

TIME SERIES

INACTIVE

- ☐ Show data download links
- ☐ Ignore outliers in chart scaling

Tooltip sorting method:

default

Smoothing

0.6

Horizontal Axis

STEP

RELATIVE

WALL

Runs

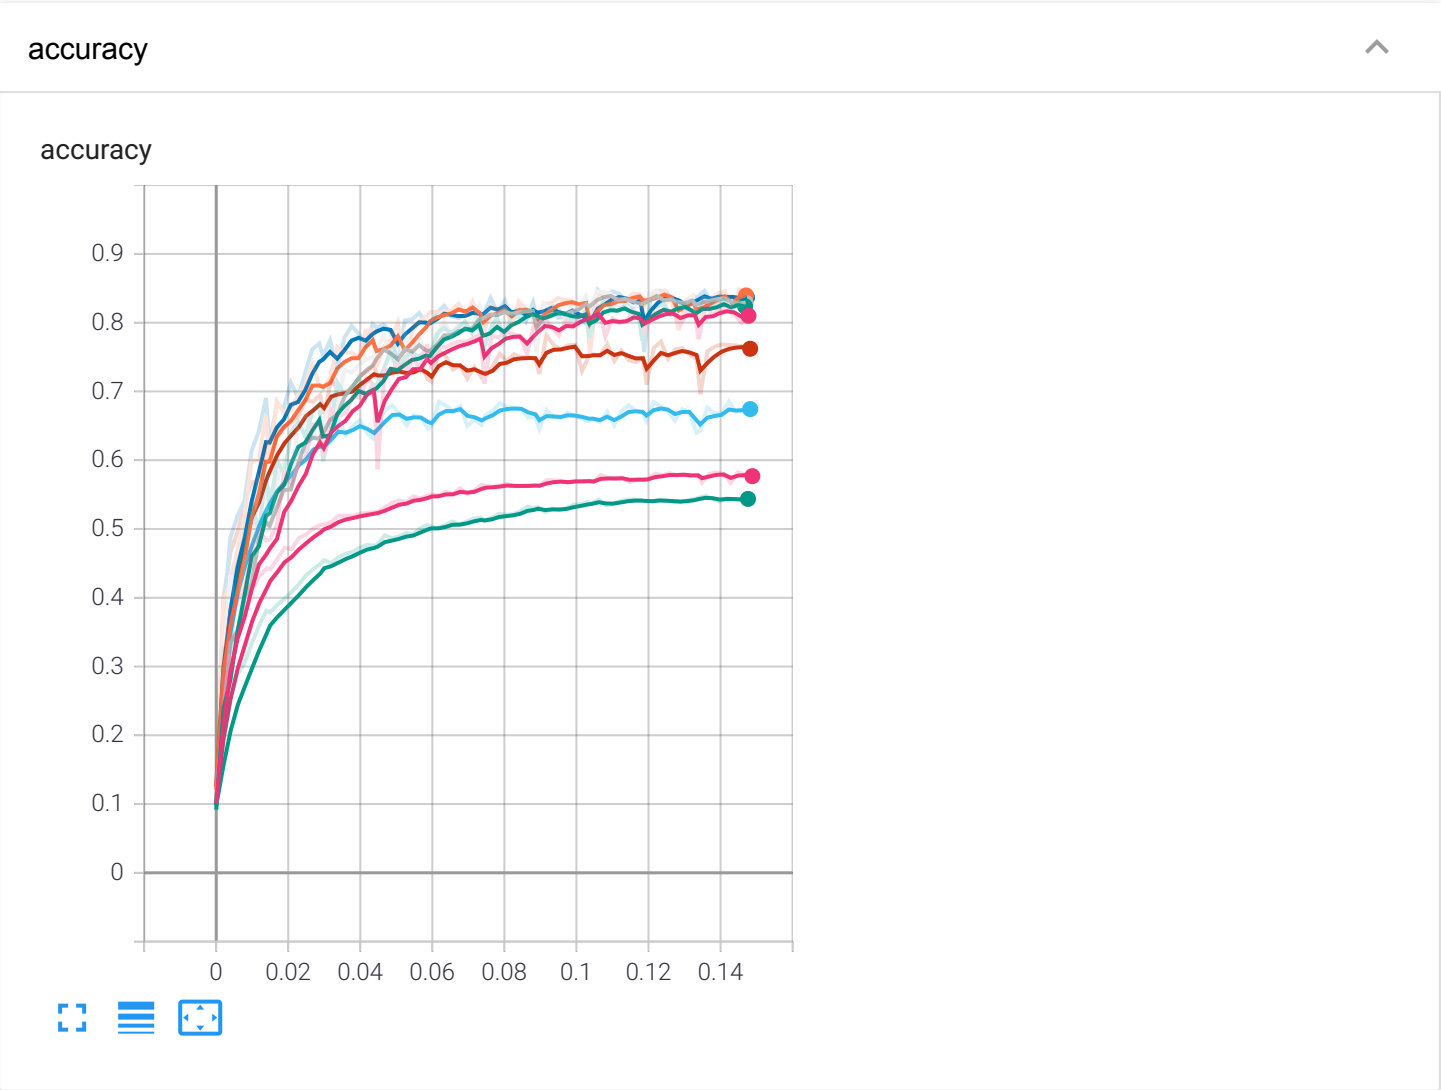
Write a regex to filter runs

- 9da8lr=(0.005,)
- Jan31\_15-52-59\_74fe5298  
9da8lr=(0.002,)
- Jan31\_16-02-49\_74fe5298  
9da8lr=(0.0015,)
- Jan31\_16-12-39\_74fe5298  
9da8lr=(0.001,)
- Jan31\_16-22-30\_74fe5298  
9da8lr=(0.0005,)
- Jan31\_16-32-22\_74fe5298  
9da8lr=(0.0001,)

TOGGLE ALL RUNS

runs

Filter tags (regular expressions supported)



Part 3

The code provided below will allow you to visualise the feature maps computed by different layers of your network. Run the code (install matplotlib if necessary) and **answer the following questions**:

1. Compare the feature maps from low-level layers to high-level layers, what do you observe?
2. Use the training log, reported test set accuracy and the feature maps, analyse the performance of your network. If you think the performance is sufficiently good, explain why; if not, what might be the problem and how can you improve the performance?
3. What are the other possible ways to analyse the performance of your network?

YOUR ANSWER FOR PART 3 HERE

A: Q1: Compare the feature maps from low-level layers to high-level layers, what do you observe?

Visualizing how a CNN learns to identify different features present in images provides a deeper insight into how the model works. It will also help to understand why the model might be failing to classify some of the images correctly and hence fine-tuning the model for better accuracy and precision. These methods aim at learning high-level semantic features via a hierarchical architecture to predict image categories.

We can see that the feature maps closer to the input of the model capture a lot of fine detail in the image, ie: edges and blobs. For the example image we have below of a truck, we can clearly see that after the convolutional layer distinctive and discriminative features such as the edges and corners of the truck light up in different feature maps. Therefore we have detected multiple features in this single convolutional layer. The next layer combines previous features to extract more complex representations as it learns features of features of the first layer. When we progress deeper into the model, the feature maps show less and less detail and the model learns higher level features as we have lower spatial resolution.

This pattern is expected, as the model abstracts the features from the image into more general concepts that can be used to make a classification. We generally lose the ability to interpret these deeper feature maps.

**Q2: Use the training log, reported test set accuracy and the feature maps, analyse the performance of your network. If you think the performance is sufficiently good, explain why; if not, what might be the problem and how can you improve the performance?**

Considering the number of epochs is only 10, I believe this training is to a high level. The training log shows that the accuracy on the validation sets increased sequentially from epoch to epoch showing the system was learning. In addition, the test accuracy of 83.39% is a very good result and it reveals a good performance of the network.

The low-level layers of the feature maps, as seen below, extract hierachical, discrimitive and distinctive features. These inlude but are not limited to edges and corners of the truck. As we move into deeper feature maps, such as layers 3 and 4, we can see that nearly all of the light ups in the feature maps are associated with the truck, showing that our network is recognising many different higher level features associated with the truck.

Even though the overall performance is good, even relative to the state of the art system only a few years ago, there are still places in the network where improvements could be implemented to increase the overall performance. The first beneficial experiment would be to vary the train / validation / test split to find the optimal split. We could introduce drop out and could vary the batchsize to see if either have a positive impact on our accuracy. There is a strong potential for success by trying out different optimizers such as adagrad, Rmsprop, adam adaddelta, adamax and nadam. Finally another potential improvement would be to vary the number of layers in the ResNet, ie: make ResNet-50.

**Q3: What are the other possible ways to analyse the performance of your network?**


Accuracy is a good measure to use here, but there are a variety of other evaluation metrics that we could use to analyse the performance of the network. Drawing out the entire confusion matrix for our data is a very good way to analyse how our data is spread out between true positives, false positives, false negatives and true negatives, and to see if we have a balanced data distribution. From the confusion matrix we can calculate the F1 score (the harmonic mean of precision and recall) and we can look at the metrics recall, precision and macro-averaged recall.

We can look at robustenss of the model in terms of the degradation in performance with respect to varying noise levels or other image artefacts.

On top of accuracy measures we want our models to be fast to train and to query, so we can look at the rate of convergence. We can analyse how scalable it is so that it will work well with large datasets, and how interpretable it is so that it is understandable and the model can explain its predictions.

```
1 torch.manual_seed(11)
2
3 data, _ = test_data[11]
4 data = data.unsqueeze_(0)
5
6 grid = torchvision.utils.make_grid(data)
7 plt.figure(figsize=(2.5,2.5))
8 plt.imshow(np.transpose(grid, (1,2,0)))
9
```

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).
<matplotlib.image.AxesImage at 0x7fcb47c382b0>




```
1 #!pip install matplotlib
2
3 torch.manual_seed(11)
4
5 import matplotlib.pyplot as plt
6
7 plt.tight_layout()
8
9
10 activation = {}
11 def get_activation(name):
12     def hook(model, input, output):
13         activation[name] = output.detach()
14     return hook
15
16 vis_labels = ['conv1', 'layer1', 'layer2', 'layer3', 'layer4']
17
18 for l in vis_labels:
19
20     getattr(model, l).register_forward_hook(get_activation(l))
21
22
23 data, _ = test_data[11]
24 data = data.unsqueeze_(0).to(device = device, dtype = dtype)
25
```

```
26 output = model(data)
27
28
29
30 for idx, l in enumerate(vis_labels):
31
32     act = activation[l].squeeze()
33
34     if idx < 2:
35         ncols = 8
36     else:
37         ncols = 32
38
39     nrows = act.size(0) // ncols
40
41     fig, axarr = plt.subplots(nrows, ncols, figsize= (32,24))
42     fig.suptitle(l)
43
44
45     for i in range(nrows):
46         for j in range(ncols):
47             axarr[i, j].imshow(act[i * nrows + j].cpu())
48             axarr[i, j].axis('off')
```




<Figure size 432x288 with 0 Axes>

conv1



layer1



1

